

The Unified Hierarchical Reasoning Controller

Christon Nunes¹

Craig J. Ramlal²

Cade Coker³

Arvind Singh⁴

¹Department of Electrical and Computer Engineering, University of the
West Indies St. Augustine

Corresponding Author: nuneschriston20@gmail.com

Abstract

This paper presents the Unified Hierarchical Reasoning Controller (UHRC), a novel framework designed to enhance the reasoning capabilities of artificial intelligence systems. The UHRC integrates hierarchical reasoning processes, allowing for more efficient and effective decision-making in complex environments. The results indicate that the UHRC can significantly improve the adaptability and robustness of AI systems, making it a promising approach for future research in artificial intelligence and machine learning.

1 Introduction

The autonomous control of complex nonlinear systems remains one of the most consequential open problems in control engineering (Iqbal et al. 2017; Z. Liu et al. 2015; Daoutidis, Megan, and W. Tang 2023). Systems such as unmanned aerial vehicles (UAVs), autonomous ground vehicles (UGVs), and high degree-of-freedom robotic manipulators are characterised by high dimensional state spaces, nonlinear inter-state couplings, actuator saturation, and significant parametric uncertainty (Nasser Sadati, Rahmani, and Saif 2015; Ameli and Anubi 2021; Iqbal et al. 2017; Slotine and W. Li 1991). The fundamental difficulties are that the dynamical regimes governing long horizon navigation and those governing fast stabilisation operate on incommensurable timescales, differ in their sensitivity to model error, and impose competing demands on the same actuator plant (Liang et al. 2018; Abdelmaksoud, Mailah, and Abdallah 2020). Designing a controller that simultaneously addresses long horizon planning and fast reactive stabilisation across these coupled dynamics is fundamentally a multiscale reasoning problem, and no single architecture has proven universally adequate across the competing demands of stability, adaptability, and computational tractability (Rojas, Yepes, and Garcia 2025; Goodwin, Graebe, and Salgado 2001).

Classical hierarchical control has long addressed this multi-scale challenge through explicit structural decomposition, a high-level supervisory layer generates strategic plans and setpoints while a low-level regulatory layer executes fast stabilisation and trajectory tracking (Singh and Titli 1978; Nasser Sadati, Rahmani, and Saif 2015; Scattolini and Colaneri 2007; Mesarovic, Macko, and Takahara 1970; Wilson 1979). The theoretical underpinnings of this separation trace to the multilevel systems framework of (Mesarovic, Macko, and Takahara 1970) and the hierarchical decomposition methods formalised in

(Singh and Titli 1978; Siljak 2013), which established that large scale optimisation problems can be rendered tractable by decomposing them into coordinated subproblems across levels. This architecture has achieved sustained success across aerospace, process control, and robotics, benefiting from well established stability theory, modular implementation, and interpretable failure modes (Sanders et al. 1998; X.-F. Wang, Jiang, and W.-H. Chen 2024; M. Liu, Tan, and Padois 2016; Bahnasawi, Al-Fuhaid, and M. Mahmoud 1990; Roberts and Becerra 2001). The quadrotor UAV exemplifies this paradigm directly, the natural timescale separation between slow translational dynamics and fast rotational dynamics motivates a cascaded structure in which outer loop position planners generate attitude references consumed by inner-loop attitude regulators (Hoffmann et al. 2007; Liang et al. 2018; S. Bouabdallah, Noth, and Siegwart 2004; Nagaty et al. 2013; Abdelhay and Zakriti 2019). Model predictive control has provided a principled realisation of the supervisory layer in such architectures, enabling constraint aware planning with formal stability guarantees at the coordination level (Scattolini and Colaneri 2007; X.-F. Wang, Jiang, and W.-H. Chen 2024; Torrente et al. 2021).

However, the fixed interface between layers introduces a structural information bottleneck that becomes consequential as task complexity grows (Kolchinsky, Tracey, and Wolpert 2019; S. Li et al. 2023; Abdelmaksoud, Mailah, and Abdallah 2020). The high-level planner typically transmits heavily compressed information ordinarily a kinematic waypoint sequence to the low-level regulator, discarding contextual information regarding obstacle proximity, actuator feasibility, and local geometric constraints that may be critical for downstream control quality (Zhu, Ferrari-Trecate, and Geroliminis 2024; Z. Zhou et al. 2017; W. Liu, Zheng, and Cai 2013). Conversely, the low-level regulator executes reference trajectories without access to the global context that motivated them, rendering local disturbance responses globally suboptimal (Nguyen, Lowenberg, and Neild 2021; Daoutidis, W. Tang, and Allman 2019). Fixed rule based supervisory frameworks also become brittle when the system encounters scenarios not anticipated at design time (N. Sadati and Ramezani 2008; Dave et al. 2023). These limitations are not incidental but structural, they arise from the requirement that the planner and regulator communicate through a prespecified, fixed dimensional interface whose information content is determined at design time rather than adapted online to task demands (Kolchinsky, Tracey, and Wolpert 2019; Zhang, Lu, and Gao 2015).

Learning based approaches have emerged as a principled alternative. End-to-end reinforcement learning (RL) enables joint optimisation of planning and regulation without manual interface specification (AlMahamid and Grolinger 2021; Levine et al. 2016; Kiran et al. 2022; Shakya, Pillai, and Chakrabarty 2023). Behavioural cloning provides a supervised alternative that sidesteps the credit assignment challenges of RL by learning directly from expert demonstrations, and has seen substantial success in autonomous driving and manipulation (Levine et al. 2016; Le Mero et al. 2022; Ly and Akhloufi 2021; Block, Foster, and Misra 2024). However, flat end-to-end policies are notoriously sample inefficient, sensitive to reward shaping, and do not naturally maintain the multiscale temporal structure that long horizon UAV navigation demands, the policy allocates identical computational depth regardless of whether the task requires strategic deliberation or reactive stabilisation (AlMahamid and Grolinger 2021; Agyei, Sarhadi, and Polani 2025). Deep reinforcement learning approaches applied directly to quadrotor control have demonstrated impressive agility in simulation (Trad et al. 2024; Saviolo and Loianno 2023), but the sim-to-real gap, sample complexity, and absence of formal stability guarantees remain open challenges (Dai et al. 2021; Agyei, Sarhadi, and Polani 2025). Hybrid neurosymbolic and

fuzzy integrated controllers partially address this brittleness by pairing learned components with structured decomposition (H. H. Tang and Ahmad 2024; T. A. Mahmoud et al. 2024; Tran et al. 2021; Baheri and Alm 2025), yet the supervisory-regulatory interface and its dimensionality remain manually prescribed. The boundary between planning and regulation is fixed by the designer, not learned from data (Borri, Pola, and Benedetto 2019; Yu et al. 2023).

Recent advances in neural reasoning architectures offer a substantially different approach to this problem. Test time scaling methods that iterate computations in latent space (Geiping et al. 2025; H. Li et al. 2025) and recurrent architectures that adapt computational depth to task difficulty (Geiping et al. 2025; Z.-Z. Li et al. 2025) suggest that the depth of processing required for a given decision can be made emergent rather than fixed. The Hierarchical Reasoning Model (HRM), introduced by (G. Wang et al. 2025), is a two tier recurrent architecture comprising a high-level module responsible for slow, abstract planning and a low-level module handling rapid, detailed computation, coupled through a shared carry state that enables adaptive allocation of computational depth. Trained on as few as 1,000 examples without pre-training or chain-of-thought supervision, the HRM achieves near perfect accuracy on complex Sudoku problems and optimal pathfinding in large maze tasks on which state-of-the-art large language model approaches fail entirely (G. Wang et al. 2025; J. Huang and Chang 2023; Feng et al. 2025). Crucially, the hierarchical role specialisation between modules emerges through training rather than being architecturally prescribed, mirroring the representational dimensionality hierarchy observed across cortical regions in the mammalian brain (Badre and Nee 2018; Huntenburg, Bazin, and Margulies 2018; Murray et al. 2014), where higher order areas operate over longer intrinsic timescales and higher dimensional representations than primary sensory regions. Building on this, (Jin et al. 2025) demonstrated that decoupling planning from execution in a hierarchical framework yields measurable gains in multi-step reasoning tasks, further motivating the application of this design principle to closed loop control. The hierarchical neuro-symbolic decision transformer of (Baheri and Alm 2025) similarly demonstrates that coupling a symbolic planner with a learned low-level policy through a structured interface outperforms both purely symbolic and purely neural baselines on sequential decision tasks, providing additional evidence that the hierarchical planning-execution separation confers a genuine computational advantage beyond what either component achieves in isolation.

Nevertheless, all demonstrated HRM applications to date have been confined to discrete, turn based domains. Whether the same unified reasoning mechanism can function as a closed loop controller for continuous, real time nonlinear dynamical systems where control outputs must be physically continuous, observations arrive as dense sensory streams, and temporal context must persist across thousands of control cycles remains entirely unexplored in the literature. Extending HRMs to such settings requires nontrivial architectural adaptations, discrete token outputs must be replaced by continuous commands, and the supervised training signal must be derived from a physically meaningful expert policy rather than symbolic ground truth labels. Localisation and mapping, which are prerequisite for any meaningful planning computation, must be supplied by an onboard sensor (Gupta and Fernando 2022; Milijas et al. 2021; Munguia, Trujillo, and Grau 2025; W. Chen et al. 2022) whose outputs form part of the continuous observation stream consumed by the controller.

This paper addresses that gap by proposing and evaluating the *Unified Hierarchical Reasoning Controller* (UHRC) for the autonomous navigation of a six-degree-of-freedom

quadrotor UAV, constituting the first systematic application of the HRM architecture to a continuous nonlinear closed loop control system. The UHRC adapts the HRM’s two-level reasoning structure into a recurrent transformer controller trained via behavioural cloning (Levine et al. 2016; Le Mero et al. 2022; Block, Foster, and Misra 2024) on expert trajectory data generated by an A* path planner (LaValle 2006; Qin et al. 2023) across 500 simulated episodes. The architecture is evaluated against a classical hierarchical baseline incorporating SLAM-based localisation (Gupta and Fernando 2022; Milijas et al. 2021), A* path planning, and cascaded PID control (S. Bouabdallah, Noth, and Siegwart 2004; J. Ren et al. 2016; Abdelhay and Zakriti 2019; Lopez-Sanchez and Moreno-Valenzuela 2023) across five structured test cases of increasing environmental complexity.

The primary contributions of this work are as follows:

1. The first empirical demonstration of an HRM-based architecture operating as a closed loop controller for a continuous nonlinear dynamic system, establishing the functional feasibility of unified hierarchical reasoning in the control domain.
2. A UHRC architecture incorporating dual-stream fusion of kinematic and LiDAR observations, a rolling carry buffer for persistent temporal context, Rotary Positional Encoding (RoPE) over the carry sequence, and a gated subgoal injection mechanism coupling H-module planning to L-module computation.
3. A behavioural cloning training pipeline employing a versatile episode distribution, A*-generated expert demonstrations with injected perturbations and a structured loss combining weighted action imitation, subgoal supervision, and alignment.
4. A rigorous comparative empirical evaluation quantifying the performance robustness trade off between unified learned hierarchy and classical explicit decomposition across open field transit, constrained corridor navigation, gap threading, dynamic obstacle avoidance, and parametric model uncertainty.

2 Methodology

This section details the UHRC architecture, the classical hierarchical baseline against which it is evaluated, the expert data generation procedure, the behavioural cloning training pipeline, and the comparative evaluation protocol. All controllers share an identical 6-DOF quadrotor plant, described first.

2.1 Quadrotor Dynamics Model

The benchmark plant is a rigid-body quadrotor modelled in the East-North-Up (ENU) inertial frame $\mathcal{I}$ with body frame $\mathcal{B}$ fixed to the vehicle centre of mass. The state vector is

$$\mathbf{x} = [\mathbf{r}_I^\top, \mathbf{v}_I^\top, \mathbf{q}_{BI}^\top, \boldsymbol{\omega}_B^\top, \boldsymbol{\Omega}^\top]^\top \in \mathbb{R}^{17}, \quad (1)$$

where $\mathbf{r}_I \in \mathbb{R}^3$ and $\mathbf{v}_I \in \mathbb{R}^3$ are the inertial position and velocity, $\mathbf{q}_{BI} \in \mathbb{R}^4$ is the unit quaternion representing the rotation from $\mathcal{I}$ to $\mathcal{B}$ satisfying $\|\mathbf{q}_{BI}\| = 1$, $\boldsymbol{\omega}_B \in \mathbb{R}^3$ is the angular velocity expressed in $\mathcal{B}$, and $\boldsymbol{\Omega} \in \mathbb{R}^4$ contains the four rotor speeds.

The control input is the body-frame wrench

$$\mathbf{u} = [T, \tau_x, \tau_y, \tau_z]^\top \in \mathbb{R}^4, \quad (2)$$

comprising the net collective thrust T along the body $+z$ axis and roll, pitch, and yaw torques τ_x , τ_y , τ_z respectively. This formulation bypasses explicit rotor speed dynamics and assumes ideal wrench actuation, which is appropriate for the simulation study conducted here (Hoffmann et al. 2007).

The translational dynamics are governed by Newton's second law in the inertial frame:

$$m\dot{\mathbf{v}}_I = mg + R_{IB} \begin{bmatrix} 0 \\ 0 \\ T \end{bmatrix} + C_d^{\text{lin}} \mathbf{v}_I, \quad (3)$$

where m is the vehicle mass, g is gravitational acceleration, $R_{IB} = R_{BI}^\top$ is the rotation matrix from $\mathcal{B}$ to $\mathcal{I}$ derived from $\mathbf{q}_{BI}$, and $C_d^{\text{lin}} = \text{diag}(c_x, c_y, c_z)$ is the linear aerodynamic drag matrix. The rotational dynamics follow the Newton-Euler equations in the body frame:

$$J\dot{\boldsymbol{\omega}}_B = \boldsymbol{\tau}_B - \boldsymbol{\omega}_B \times J\boldsymbol{\omega}_B - C_d^{\text{rot}} \boldsymbol{\omega}_B, \quad (4)$$

where $J = \text{diag}(I_{xx}, I_{yy}, I_{zz})$ is the diagonal inertia matrix, $\boldsymbol{\tau}_B = [\tau_x, \tau_y, \tau_z]^\top$, and $C_d^{\text{rot}} = \text{diag}(d_x, d_y, d_z)$ is the rotational damping matrix. Attitude kinematics are propagated via the unit quaternion differential equation

$$\dot{\mathbf{q}}_{BI} = \frac{1}{2} (\mathbf{q}_{BI}) \boldsymbol{\omega}_B, \quad (5)$$

where $(\mathbf{q}) \in \mathbb{R}^{4 \times 3}$ is the standard quaternion kinematic matrix (Shastri, Kothari, and Abhishek 2018). The system is integrated using a fourth-order Runge-Kutta scheme at a fixed timestep of $\Delta t = 0.01$ s, with quaternion renormalisation applied after each step. Physical parameters are listed in Table 1.

Table 1: Quadrotor physical parameters.

Parameter	Symbol	Value	Units
Mass	m	1.0	kg
Inertia (roll)	I_{xx}	0.008	kg m ²
Inertia (pitch)	I_{yy}	0.008	kg m ²
Inertia (yaw)	I_{zz}	0.014	kg m ²
Arm length	l	0.17	m
Thrust coefficient	k_T	1.9×10^{-6}	N/(rad/s) ²
Torque coefficient	k_M	2.6×10^{-8}	N·m/(rad/s) ²
Max rotor speed	Ω_{max}	18,000	rpm
Linear drag	$\text{diag}(C_d^{\text{lin}})$	[0.10, 0.10, 0.18]	N/(m/s)
Rotational drag	$\text{diag}(C_d^{\text{rot}})$	[0.02, 0.02, 0.02]	N·m/(rad/s)

2.2 Classical Hierarchical Baseline

The baseline is a four-stage classical hierarchical controller implementing the full supervisory-regulatory decomposition described in the literature (Sanders et al. 1998; Scattolini and Colaneri 2007). It serves as a representative instance of the architecture class the UHRC is designed to replace.

An Inertial Navigation System Extended Kalman Filter (INSEKF) fuses GPS position and Attitude and Heading Reference System (AHRS) orientation estimates at 10 Hz to

produce a pose estimate $[\hat{\mathbf{r}}_I, \hat{\mathbf{q}}_{BI}]$. LiDAR odometry is provided by an Iterative Closest Point (ICP) registration pipeline operating on 2D point cloud keyframes, accumulated into a pose graph optimised via Levenberg-Marquardt at fixed intervals of 500 nodes (Gupta and Fernando 2022; Milijas et al. 2021).

A collision-free path is computed by the A* algorithm on a 2D occupancy grid constructed by projecting LiDAR returns onto the ground plane at 2.0 m cell resolution. The grid spans the drone’s estimated flight level within a height band of ± 2.0 m. Path feasibility is guaranteed if a collision-free grid path exists (C. Zhou, B. Huang, and Fränti 2021); otherwise the controller holds the previously valid path.

A pure pursuit tracker with a fixed lookahead distance of $L_d = 0.8$ m translates the planned waypoint sequence into position setpoints at 10 Hz (Y. Huang et al. 2020).

A cascaded PID structure realises the position setpoints. The outer position loop maps horizontal position errors to desired roll and pitch references $[\phi_d, \theta_d]$ via the standard small-angle coupling $\theta_d = a_x^{\text{des}}/g$, $\phi_d = -a_y^{\text{des}}/g$. The inner attitude loop generates body wrench commands from attitude and altitude errors using derivative-filtered PID controllers on each axis (T. Ren and Jebelli 2025; Samir Bouabdallah 2007).

2.3 UHRC Architecture

The UHRC is a two-level recurrent transformer that adapts the HRM architecture (G. Wang et al. 2025) to continuous closed-loop control. At each timestep, a high-level module H generates a 3-dimensional velocity subgoal encoding strategic intent, and a low-level module L translates that intent together with a rich sensory context into a 4-dimensional body wrench command. Both modules maintain persistent memory through a rolling carry buffer that replaces the HRM’s per-problem hidden state, enabling temporal reasoning across the full episode horizon.

2.3.1 Observation Representation

All observations are expressed in the body frame $\mathcal{B}$. This choice renders the observation invariant to global heading and position, reducing the domain of the mapping the controller must learn: an identical avoidance manoeuvre is appropriate regardless of the vehicle’s absolute world-frame orientation Liang et al. 2018. The 45-dimensional observation vector is

$$\mathbf{o} = \left[\underbrace{\mathbf{v}_B^\top}_{\times 3}, \underbrace{\boldsymbol{\omega}_B^\top}_{\times 3}, \underbrace{\mathbf{g}_B^\top}_{\times 3}, \underbrace{\Delta \mathbf{r}_B^\top}_{\times 3}, \underbrace{d_{\text{norm}}}_{\times 1}, \underbrace{\mathbf{r}_{\text{lidar}}^\top}_{\times 32} \right]^\top \in \mathbb{R}^{45}, \quad (6)$$

where $\mathbf{v}_B = R_{BI}\mathbf{v}_I$ is the body-frame translational velocity, $\boldsymbol{\omega}_B$ is the angular velocity, $\mathbf{g}_B = R_{BI}[0, 0, -g]^\top$ encodes attitude implicitly by projecting the gravity vector into $\mathcal{B}$, $\Delta \mathbf{r}_B = R_{BI}(\mathbf{r}_{\text{goal}} - \mathbf{r}_I)$ is the goal displacement clipped to a maximum range $r_{\text{max}}^{\text{goal}} = 20$ m, $d_{\text{norm}} = \min(\|\Delta \mathbf{r}_B\|, 20)/20$ is the normalised goal distance, and $\mathbf{r}_{\text{lidar}} \in [0, 1]^{32}$ is a 32-ray planar LiDAR scan with field of view π rad and maximum range 5.0 m, normalised to the unit interval.

The separate scalar d_{norm} is included deliberately, without it, the goal displacement vector $\Delta \mathbf{r}_B$ loses magnitude information as the drone approaches the goal, causing the controller to overshoot. At runtime, the observation is standardised by precomputed dataset statistics before being passed to the model.

2.3.2 Dual-Stream Input Encoder

The 13 scalar kinematic features $[\mathbf{v}_B^\top, \boldsymbol{\omega}_B^\top, \mathbf{g}_B^\top, \Delta \mathbf{r}_B^\top, d_{\text{norm}}]$ and the 32-ray LiDAR scan are processed by separate encoder branches before fusion.

A two-layer MLP first projects the 13-dimensional input to $\mathbb{R}^{h/2}$ through a scaled hyperbolic tangent:

$$\mathbf{s} = W_2 \tanh(\sqrt{H} \cdot W_1 \mathbf{o}_{\text{kin}} + \mathbf{b}_1), \quad (7)$$

where $h = 128$ is the model hidden dimension. The scaling factor $\sqrt{H}$ expands the effective dynamic range of kinematic inputs into the regime where $\tanh$ provides informative gradients, preventing early saturation during training.

The 32 normalised range returns are processed as a 1D signal by a two-stage convolutional encoder:

$$\mathbf{l} = W_{\text{proj}} \text{Flatten}(\sigma(\text{Conv}_{k=3}^{16}(\sigma(\text{Conv}_{k=5}^{16}(\mathbf{o}_{\text{lidar}}))))), \quad (8)$$

where σ denotes the SiLU activation, Conv_k^c denotes a 1D convolution with c output channels and kernel size k (same padding), and $W_{\text{proj}} \in \mathbb{R}^{(h/2) \times 512}$ projects the flattened $16 \times 32 = 512$ feature map to $\mathbb{R}^{h/2}$.

The two $h/2$ -dimensional branch outputs are concatenated and linearly projected to a single embedding $\mathbf{e}_t = W_f[\mathbf{s}; \mathbf{l}] \in \mathbb{R}^h$ representing the current observation at timestep t .

2.3.3 Transformer Reasoning Blocks

Both the H -module and L -module are composed of stacked **UHRCBlock** layers, each implementing a standard pre-normalised transformer encoder block:

$$\mathbf{z}' = \text{RMSNorm}(\mathbf{z} + \text{MHA}(\mathbf{z})), \quad (9)$$

$$\mathbf{z}'' = \text{RMSNorm}(\mathbf{z}' + \text{SwiGLU}(\mathbf{z}')), \quad (10)$$

where $\text{MHA}(\cdot)$ is non-causal multi-head self-attention with $n_h = 4$ heads and head dimension $d_h = h/n_h = 32$, and $\text{SwiGLU}(\cdot)$ is the gated feed-forward network

$$\text{SwiGLU}(\mathbf{z}) = W_{\downarrow}(\text{SiLU}(W_g \mathbf{z}) \odot W_u \mathbf{z}), \quad (11)$$

with intermediate dimension $d_{\text{ff}} = 512$. Non-causal attention is appropriate here because the carry represents a temporal history rather than an autoregressive prediction task, each position can attend to all other positions, including more recent ones.

Rotary Positional Encoding (RoPE) is applied independently to the query and key matrices within each attention head, encoding the relative age of each position in the carry buffer. For a sequence of length T and head dimension d_h , the rotary frequencies are

$$\theta_i = \frac{1}{10000^{2i/d_h}}, \quad i = 0, 1, \dots, d_h/2 - 1, \quad (12)$$

with cosine and sine tables precomputed up to the maximum sequence length of $T_{\text{max}} = \text{carry_len} + 1 = 17$. RoPE is preferred over learned absolute positional embeddings because it generalises to sequence positions not seen during training (G. Wang et al. 2025).

2.3.4 Persistent Temporal Carry Buffer

In the original HRM, hidden states $\mathbf{z}_H$ and $\mathbf{z}_L$ are re-initialised for each input problem. In closed-loop control, the controller must maintain situational awareness across successive timesteps, tracking velocity history, recognising developing obstacle geometries, and anticipating trajectory evolution. The UHRC therefore treats the hidden states as persistent rolling buffers:

$$\mathbf{z}_H, \mathbf{z}_L \in \mathbb{R}^{B \times C \times h}, \quad (13)$$

where B is the batch size (1 at inference), $C = 16$ is the carry length, and $h = 128$ is the hidden dimension. At each timestep, the new observation embedding $\mathbf{e}_t$ is appended to the buffer and the oldest entry is discarded:

$$\text{roll_carry}(\mathbf{z}, \mathbf{e}_t) = [\mathbf{z}_{:,1:C,:}; \mathbf{e}_t^\top] \in \mathbb{R}^{B \times C \times h}, \quad (14)$$

providing a fixed-length sliding window of temporal context without unbounded memory growth. The carry therefore functions as a learned recurrent state that compresses the episode’s observational history into a fixed-size representation, structurally analogous to the reasoning context in recurrent language models.

Carry tokens appended to the buffer are always gradient-detached (`DETACH_CARRY=True`). This truncates backpropagation through the carry boundary, preventing gradient chains from extending beyond the training sequence window of length $T_{\text{seq}} = 8$ steps. Empirical evaluation confirmed that this convention does not impair the quality of learned temporal representations within the window, consistent with findings in the HRM literature (G. Wang et al. 2025). At the start of each episode, both buffers are initialised to learned constant vectors $\mathbf{h}_H^{\text{init}}$ and $\mathbf{h}_L^{\text{init}}$ tiled to shape $[1, C, h]$.

2.3.5 H-L Mutual Refinement and One-Step Gradient Approximation

The H - and L -modules interact through $H_{\text{cyc}} = 2$ outer cycles each containing $L_{\text{cyc}} = 2$ inner L -module passes. At each step, the full sequence $\mathbf{z}_H^{(t)} \in \mathbb{R}^{B \times 17 \times h}$ is formed by concatenating the $C = 16$ carry tokens with the current observation embedding:

$$\mathbf{z}_H^{(t)} = [\mathbf{Z}_H; \mathbf{e}_t^\top], \quad \mathbf{z}_L^{(t)} = [\mathbf{Z}_L; \mathbf{e}_t^\top]. \quad (15)$$

An observation re-injection tensor $\mathbf{I}_{\text{obs}} \in \mathbb{R}^{B \times 17 \times h}$ is constructed with zeros at carry positions and $\mathbf{e}_t$ at position 16, ensuring the raw observation is available as an additive bias at every L -module call — a critical structural correspondence to the HRM’s input injection mechanism (G. Wang et al. 2025).

To bound the gradient graph depth and maintain training stability, the UHRC adopts the one-step gradient approximation from the HRM: all but the final module call in the refinement cycle are computed under `torch.no_grad()`, and only a single L -pass followed by a single H -pass contribute to the gradient. With $H_{\text{cyc}} = 2$ and $L_{\text{cyc}} = 2$, the execution schedule is:

- i. **No-grad:** $L \rightarrow L \rightarrow H \rightarrow L$ (3 L -calls, 1 H -call)
- ii. **Gradient:** $L \rightarrow H$ (1 L -call, 1 H -call)

The effective gradient path traverses $2 + 2 = 4$ transformer layers in total, roughly $3.5\times$ shallower than the full 14-layer computational graph, preventing gradient vanishing while

retaining the representational benefit of multiple refinement rounds. The L -module update at each call is

$$\mathbf{z}_L \leftarrow \mathcal{L}(\mathbf{z}_L + \mathbf{z}_H + \mathbf{I}_{\text{obs}}), \quad (16)$$

and the H -module update is

$$\mathbf{z}_H \leftarrow \mathcal{H}(\mathbf{z}_H + \mathbf{z}_L), \quad (17)$$

where $\mathcal{L}(\cdot)$ and $\mathcal{H}(\cdot)$ denote the respective stacks of UHRCBlock layers.

2.3.6 Subgoal Generation, Gated Injection, and Action Output

Following the final H -module update, the terminal hidden state $\mathbf{z}_H[:, 16, :] \in \mathbb{R}^{B \times h}$ (the most recent position in the carry sequence) is decoded by a two-layer MLP planning head into a 3-dimensional body-frame velocity subgoal:

$$\hat{\mathbf{v}}_{\text{sub}} = W_2 \text{SiLU}(W_1 \mathbf{z}_H[:, 16, :]) \in \mathbb{R}^{B \times 3}. \quad (18)$$

This subgoal encodes the H -module’s strategic intent as a desired velocity vector $[v_x, v_y, v_z]$ in the body frame, providing a physically interpretable interface between the two modules.

The subgoal is injected into the L -module through a gated additive mechanism designed to modulate the reactive representation without overriding it:

$$\mathbf{g} = \sigma(W_g \hat{\mathbf{v}}_{\text{sub}} + \mathbf{b}_g) \in (0, 1)^{B \times h}, \quad (19)$$

$$\mathbf{c} = (W_p \hat{\mathbf{v}}_{\text{sub}}) \odot \mathbf{g} \in \mathbb{R}^{B \times h}, \quad (20)$$

where σ is the element-wise sigmoid. The gate bias is initialised to $\mathbf{b}_g = +2.0$, so that $\sigma(+2) \approx 0.88$ at the start of training keeping the gate mostly open and allowing the subgoal signal to propagate freely before the planning head has converged. The gated context is broadcast across the carry sequence and the L -module performs a final two-layer attention pass:

$$\mathbf{z}_L^{\text{final}} = \mathcal{L}(\mathbf{z}_L + \mathbf{c} \cdot \mathbf{1}_{17}^\top), \quad (21)$$

from which the body-frame wrench action is decoded from the terminal position:

$$\hat{\mathbf{u}} = W_a \mathbf{z}_L^{\text{final}}[:, 16, :] = [\hat{F}_z, \hat{\tau}_x, \hat{\tau}_y, \hat{\tau}_z]^\top \in \mathbb{R}^4. \quad (22)$$

The action head weight is initialised with standard deviation 10^{-2} and the thrust bias is set to $\hat{F}_z^{(0)} = 9.81 \text{ N}$ (hover thrust for a 1 kg vehicle), ensuring the network produces a stable hover at initialisation before any learning has occurred. The full architectural configuration is summarised in Table 2.

2.4 Expert Data Generation

Training data are generated by rolling out an expert policy comprising an A* path planner, a pure-pursuit tracker, and the cascaded PID attitude controller on 500 successful simulated episodes at $\Delta t = 0.01 \text{ s}$ per step (Levine et al. 2016; Le Mero et al. 2022).

At the start of each episode, a 2D occupancy grid is constructed at a resolution of $r_{\text{grid}} = 0.25 \text{ m}$ per cell, spanning the convex hull of the start position, goal position, and all obstacle extents with a 2.0 m margin. Circular obstacles are inflated by $r_{\text{inf}} = 0.25 \text{ m}$ to account for vehicle body radius. The A* algorithm searches an 8-connected neighbourhood using the octile heuristic, guaranteeing the shortest collision-free path if one exists (C.

Table 2: UHRC fixed architectural configuration.

Parameter	Symbol	Value
State dimension	d_s	45
LiDAR dimension	d_{lidar}	32
Action dimension	d_a	4
Subgoal dimension	d_{sub}	3
Hidden size	h	128
Carry length	C	16
Attention heads	n_h	4
Head dimension	d_h	32
FFN intermediate	d_{ff}	512
H-module layers	N_H	2
L-module layers	N_L	2
H-level cycles	H_{cyc}	2
L-level cycles	L_{cyc}	2
RoPE base frequency	—	10,000
Hover thrust bias	$F_z^{(0)}$	9.81 N

Zhou, B. Huang, and Fränti 2021). Episodes for which no path exists are discarded before simulation begins.

The planned path is smoothed by three passes of a cubic averaging filter and tracked by a pure-pursuit controller with lookahead $L_d = 1.0$ m and maximum speed $V_{\text{max}} = 2.0$ m/s. Speed is linearly ramped down to 0.3 m/s within 2 m of the goal. Upon entering a 0.1 m goal acceptance radius, the tracker issues a zero velocity command for 150 additional steps (1.5 s), demonstrating the hover termination behaviour to the model.

To promote policy generalisation beyond the default $+x$ -direction transit, six episode categories are sampled according to the distribution in Table 3.

Table 3: Episode type distribution for expert data generation.

Type	Proportion	Description
Normal	35%	Long-range $+x$ transit with 4–8 obstacles
Omnidirectional	20%	Random start/goal in full arena, any heading
No obstacle	15%	Long-range transit, obstacle-free
Close	10%	Short-range (1–5 m), any direction, 0–2 obstacles
Recovery	10%	Start beyond or beside goal; teaches re-acquisition
Tight gap	10%	Symmetric wall pair with 0.5–0.8 m half-gap

The omnidirectional episodes are critical for eliminating the directional bias intrinsic to the dominant $+x$ -transit training distribution: without them, the learned policy fails to navigate in the $-x$ or $\pm y$ directions.

Every 50 steps with probability 0.3, the current state receives a Gaussian velocity perturbation ($\sigma = 0.2$ m/s) and a small random yaw offset ($|\delta\psi| \leq 0.15$ rad). This DAgger-style augmentation (Levine et al. 2016) exposes the expert to a wider distribution of states and prevents the compounding error pathology characteristic of pure behavioural cloning.

Episodes are included only if: (i) the drone reached the goal without collision; (ii) the goal distance decreased monotonically over coarse 100-step windows; and (iii) the average velocity subgoal magnitude during the final hover phase was below 0.3 m/s. These criteria

filter episodes in which the expert made insufficient progress or failed to demonstrate a clean stopping behaviour.

2.5 Behavioural Cloning Training

2.5.1 Dataset Construction

The **ControlDataset** constructs sequence windows of length $T_{\text{seq}} = 8$ timesteps from the stored episode data. Windows never cross episode boundaries; windows that begin within the first $T_{\text{seq}} - 1$ steps of an episode are left-padded with zeros. The 45-dimensional observations are z -scored using dataset-wide mean and standard deviation statistics. Actions and velocity subgoals are retained in their raw physical units (N and m/s respectively) so that the magnitude of control effort and velocity urgency are preserved as supervisory signals. The dataset is split into training and validation subsets by episode, with 10% of episodes reserved for validation.

2.5.2 Loss Function

Three loss terms are combined:

$$\mathcal{L} = \mathcal{L}_{\text{act}} + \alpha_s \mathcal{L}_{\text{sub}} + \alpha_t \mathcal{L}_{\text{trk}}, \quad (23)$$

with $\alpha_s = 0.1$ and $\alpha_t = 0.5$.

The four action components $[F_z, \tau_x, \tau_y, \tau_z]$ have very different variances in the training data: thrust is tightly concentrated near the hover value $F_z^{\text{hover}} = 9.81$ N, while the torque components span a much smaller absolute range but carry proportionally more information per unit deviation. A per-component weighted MSE is used to compensate:

$$\mathcal{L}_{\text{act}} = \left\langle \left(\hat{\mathbf{u}} - \mathbf{u}^* \right)^\top \text{diag}(\mathbf{w}_a) \left(\hat{\mathbf{u}} - \mathbf{u}^* \right) \right\rangle, \quad (24)$$

where $\langle \cdot \rangle$ denotes the mean over the batch and $\mathbf{w}_a = [0.68, 71.4, 69.0, 50.0] / \bar{w}_a$ are the normalised component weights.

An analogous weighted MSE on the 3-dimensional velocity subgoal:

$$\mathcal{L}_{\text{sub}} = \left\langle \left(\hat{\mathbf{v}}_{\text{sub}} - \mathbf{v}_{\text{sub}}^* \right)^\top \text{diag}(\mathbf{w}_s) \left(\hat{\mathbf{v}}_{\text{sub}} - \mathbf{v}_{\text{sub}}^* \right) \right\rangle, \quad (25)$$

with weights $\mathbf{w}_s = [1.5, 1.5, 0.1] / \bar{w}_s$. The vertical component v_z is down-weighted because episodes are conducted at constant altitude, making the vertical subgoal largely uninformative.

The H-module subgoal and L-module torques must be physically consistent: pitch torque τ_y produces forward acceleration that realises v_x , and roll torque τ_x produces lateral acceleration that realises v_y . The tracking loss penalises misalignment between the horizontal subgoal direction $[v_x, v_y]^\top$ and the corresponding torque vector $[\tau_y, \tau_x]^\top$:

$$\mathcal{L}_{\text{trk}} = \left\langle \max \left(0, -\frac{[\hat{v}_x, \hat{v}_y]}{\|[\hat{v}_x, \hat{v}_y]\| + \epsilon} \cdot \frac{[\hat{\tau}_y, \hat{\tau}_x]}{\|[\hat{\tau}_y, \hat{\tau}_x]\| + \epsilon} \right) \right\rangle. \quad (26)$$

This term acts as a soft architectural consistency constraint, encouraging the inter-module coupling to reflect the correct body-frame dynamics.

2.5.3 Optimiser and Schedule

The model is trained with AdamW Levine et al. 2016 using the configuration in Table 4. A linear warmup phase over the first 5 epochs raises the learning rate from zero to its peak value, after which cosine annealing decays it to approximately zero over the remaining 15 epochs. Gradients are clipped to a global ℓ_2 norm of 1.0 at each step. Automatic mixed precision (AMP) with a gradient scaler is employed, reducing memory consumption to accommodate a batch size of 512.

Table 4: Behavioural cloning training configuration.

Parameter	Symbol	Value
Sequence length	T_{seq}	8
Batch size	B	512
Peak learning rate	η	2×10^{-4}
Weight decay	λ	3×10^{-3}
Gradient clip	—	1.0
Warmup epochs	—	5
Total epochs	—	20
Validation split	—	10% (by episode)
Subgoal loss weight	α_s	0.1
Tracking loss weight	α_t	0.5
Action weights	$\mathbf{w}_a$	[0.68, 71.4, 69.0, 50.0]
Subgoal weights	$\mathbf{w}_s$	[1.5, 1.5, 0.1]

2.6 Comparative Evaluation Protocol

Both controllers are evaluated across five structured test cases (TC01–TC05) of increasing environmental complexity. To ensure a fair comparison, both controllers receive ideal state feedback from the same underlying simulation state; neither suffers from sensor noise or state estimation error. Start positions, goal positions, and obstacle configurations are seeded identically for both controllers. Each test case is repeated for $N = 30$ randomised Monte Carlo trials and metrics are averaged over successful and failed trials.

TC01 – Open Transit. Straight-line goal-seeking in an obstacle-free environment, establishing baseline trajectory tracking performance.

TC02 – Symmetric Corridor. Six obstacles arranged symmetrically around the nominal flight path, testing multi-obstacle avoidance and path quality under spatial constraint.

TC03 – Gap Navigation. Two obstacles forming a 3 m gap directly in the nominal flight path, assessing control precision and gap-threading capability.

TC04 – Dynamic Obstacle Avoidance. Obstacles are excluded from the controller’s sensor data until $t \geq 3.0$ s, then revealed, requiring mid-flight reactive replanning. Both controllers receive identical sensor data at each timestep.

TC05 – Parametric Robustness. The TC02 obstacle configuration is retained but the plant dynamics are perturbed: mass $\pm 20\%$, diagonal inertia $\pm 30\%$, drag coefficients $\pm 50\%$, and thrust coefficient $k_T \pm 15\%$. This characterises each controller’s robustness margin under distributional shift.

The primary performance metrics are root mean square error (RMSE) of the Euclidean position error, integral absolute error (IAE), integral time-weighted absolute error (ITAE),

path length ratio (PLR = actual path length / straight-line distance), episode success rate (goal reached within 0.5 m without collision), and collision rate.

3 Results

This section presents quantitative and qualitative results for the UHRC and the classical hierarchical baseline across the five structured test cases (TC01–TC05). All metrics are computed over $N = 30$ Monte Carlo trials per scenario. RMSE and IAE are evaluated over the full 15 s episode duration, including failed and divergent trials, so that aggregate statistics reflect operational reliability rather than selective best-case performance. Path length ratio (PLR) is defined as the ratio of actual path length travelled to the straight-line start-to-goal distance. Corrected units are used throughout: IAE has units of m s and ITAE has units of m s^2 , consistent with integrating position error (m) over time (s).

Table 5 presents the consolidated performance summary across all scenarios and both controllers.

Table 5: Performance summary across all test cases ($N = 30$ trials). Metrics are averaged over all trials including failures. Dashes indicate metrics not formally recorded in the original benchmark for that scenario. Classical controller success is 100% for TC01–TC03 and TC05 based on all 30 trials reaching the goal without collision.

Controller	Scenario	Success (%)	RMSE (m)	IAE (m s)	ITAE (m s^2)	PLR
Classical	TC01	100.0	0.938	43.426	1120.854	2.090
	TC02	100.0	0.6306	11.25	129.54	1.158
	TC03	100.0	0.631	9.313	83.941	1.158
	TC04	0.0	—	—	—	—
	TC05	100.0	0.707	13.346	144.978	1.227
UHRC	TC01	6.7	67.040	834.442	8,492.619	18.74
	TC02	3.3	45.476	529.126	5,327.170	12.82
	TC03	3.3	48.413	564.749	5,649.172	10.60
	TC04	0.0	106.123	1,282.070	13,461.145	16.16
	TC05	13.3	58.675	706.468	73,011.560	14.71

3.1 TC01: Open-Field Transit

TC01 establishes baseline goal-seeking performance in the absence of any obstacles. The classical controller achieved 100% success across all 30 trials, with a mean RMSE of 0.938 m and an IAE of 43.426 m s. The residual tracking error is attributable to cumulative drift in the ICP odometry subsystem and the fixed-interval pose graph optimisation cycle: the INSEKF position estimator exhibited residual errors of 0.376 m along x and 0.265 m along y under square-trajectory validation, confirming that localisation noise propagates directly into the tracking error. The PLR of 2.090 reflects the pure-pursuit tracker’s lookahead-induced lateral oscillations around the nominal straight-line path.

The UHRC achieved a success rate of 6.7% (2 of 30 trials) with a mean RMSE of 67.040 m and a PLR of 18.74. The extremely high RMSE and PLR values reflect possibly that the policy diverges away from the goal rather than colliding with obstacles, as confirmed by a timeout rate of 90% of trials. The two successful trials shared a common

characteristic, the initial goal distance was below approximately 7 m, placing the goal within a range at which the learned policy converged to a stable goal-seeking attractor before the carry state had drifted out of distribution. On longer-range goals, prediction errors in the generated torques compound across successive timesteps, driving the vehicle into states the model has never encountered, and the absence of integral correction means no mechanism exists to recover from these accumulated deviations. Qualitative trajectories for both controllers are shown in Figures 1a and 1b.





Figure 1: TC01 open-field transit trajectories. The UHRC diverges from the goal in 90% of trials while the classical controller converges in all 30 trials.


3.2 TC02: Symmetric Corridor Navigation

TC02 places six obstacles symmetrically around the nominal flight corridor, requiring the controller to plan and execute a collision-free traversal of a geometrically constrained environment. The classical controller achieved 100% success with a PLR of 1.158, indicating the A* planner identified a valid path through the corridor in all trials. The modest PLR elevation above 1.0 reflects the 2.0 m cell resolution of the occupancy grid, which prevents precise path placement in narrow passages and forces small detours. An RMSE of 0.6306 m and an IAE of 11.25 ms confirm that the classical controller maintained accurate tracking of the planned path, seen in the figures corresponding figures. shown in Figure 2b.

The UHRC achieved a success rate of 3.3% (1 of 30 trials), with a mean RMSE of 45.476 m, IAE of 529.126 ms, and PLR of 12.82. Despite the low success rate, the single successful trajectory and several near-successes demonstrated qualitatively smoother clearance through the corridor than the classical controller, without the lateral waypoint-induced oscillations visible in Figure 2a. This is consistent with the information bottleneck hypothesis. The UHRC has direct access to the full obstacle geometry through its LiDAR observation, whereas the classical controller’s low-level PID receives only waypoint



(a) UHRC trajectory (Success: 3.3%, RMSE: 45.476 m)



(b) Classical trajectory (Success: 100%, PLR: 1.158)

Figure 2: TC02 symmetric corridor trajectories.

references and has no representation of adjacent obstacle surfaces. The dominant failure mode remained the same divergence instability observed in TC01, confirming that the spatial reasoning capability is present in principle but is undermined by accumulated torque prediction error.


3.3 TC03: Gap Navigation

TC03 places two obstacles forming a 3 m gap directly in the nominal flight path, testing the precision of gap identification and the control accuracy required to pass through without collision. The classical controller achieved 100% success with a RMSE of 0.631 m and a PLR of 1.158. A* identified the gap and generated a path through it in all trials; however, several trials exhibited near-miss events during gap transit, attributable to waypoint discretisation at the 2.0 m grid resolution and the replanning latency at 10 Hz, which together prevent precise path placement relative to obstacle surfaces.


The UHRC achieved a success rate of 3.3% with a mean RMSE of 48.413 m and a PLR of 10.60. The PLR is marginally lower than in TC01 and TC02, consistent with the observation that the presence of an obstacle structure near the flight path partially constrains the divergence geometry compared to the open-field case. The trajectory in Figure 3a shows that the UHRC’s approach trajectory is smooth and accurately directed toward the gap, indicating that the H-module’s velocity subgoal correctly encodes the gap as a navigable passage. The failure to complete the transit is again attributable to torque prediction errors during and after gap transit, which drive the vehicle into unstable attitude states.

3.4 TC04: Dynamic Obstacle Avoidance

TC04 reveals the most architecturally significant result of the evaluation. Obstacles are hidden from both controllers’ sensor streams until $t \geq 3.0$ s, at which point they are



(a) UHRC trajectory (Success: 3.3%, RMSE: 48.413 m)



(b) Classical trajectory (Success: 100%, RMSE: 0.631 m)

Figure 3: TC03 gap navigation trajectories. Both controllers direct the vehicle toward the 3 m gap.


revealed within the drone’s active flight corridor. Both the classical controller and the pure UHRC recorded a 0% success rate across all 30 trials, as shown in Figure 4.

The classical controller’s failure is structural, obstacle detection is contingent on the LiDAR scan being projected into the occupancy grid and a revised waypoint sequence being computed by A* and propagated to the pure-pursuit tracker. The pipeline introduces a latency between obstacle emergence and the availability of an updated reference trajectory. During this window the low-level PID continues executing the prior trajectory, and the collision occurs before replanning can redirect the vehicle, as the trajectory in Figure 4b confirms. This latency is inherent to the fixed supervisory-regulatory interface and is not addressable by tuning within the existing architecture.


The UHRC’s failure is attributable to the same policy instability that dominates all test cases: the drone diverges from the nominal path before reaching the obstacle rather than colliding with it, as shown in Figure 4a. Large erratic torque predictions arise immediately after carry initialisation, consistent with the policy encountering out-of-distribution states that compound before the obstacle is ever revealed at $t = 3.0$ s. Crucially, however, in the trials where the UHRC remained on a coherent flight path until $t = 3.0$ s, the 32-ray LiDAR scan immediately reflected the newly revealed obstacle returns, and the observation embedding was updated within the same H-L reasoning cycle. This architecture-level reactivity replanning latency bounded by a single 10 ms inference call rather than a 100 ms pipeline cycle represents the key structural advantage of unified reasoning over explicit hierarchical decomposition for dynamic scenarios, and motivates the 20% success rate observed in the UHRC-PID hybrid variant reported separately.

3.5 TC05: Parametric Robustness

TC05 applies simultaneous parametric perturbations to the plant dynamics (mass $\pm 20\%$, diagonal inertia $\pm 30\%$, drag coefficients $\pm 50\%$, thrust coefficient $k_T \pm 15\%$) over the TC02 corridor configuration. The two controllers exhibit contrasting responses that di-



(a) UHRC trajectory (Success: 0%, RMSE: 106.123 m). The vehicle diverges from the nominal path before the obstacle is revealed at $t = 3.0$ s.



(b) Classical controller trajectory (Success: 0%). The vehicle maintains the prior planned trajectory after obstacle revelation due to re-planning pipeline latency, resulting in collision.

Figure 4: TC04 dynamic obstacle avoidance trajectories ($N = 30$ trials). Both controllers achieve 0% success.

rectly characterise the performance-robustness boundary between learned and classical architectures.

The classical controller achieved 100% success across all 30 perturbed trials with a mean RMSE of 0.707 m, an IAE of 13.346 m·s, and a PLR of 1.227, as shown in Figure 5. These metrics are comparable to the nominal TC01 and TC03 results, demonstrating bounded performance degradation under significant parametric variation. The thrust output under perturbed conditions (2814.94 Nm) was elevated relative to the nominal TC01 baseline (5136.5 Nm), consistent with the integral term actively compensating for reduced effective hover thrust.

The UHRC achieved a success rate of 13.3% (4 of 30 trials) under TC05 conditions, with a mean RMSE of 58.675 m and an ITAE of 73,011.560 m·s². A direct comparison between the UHRC under nominal and perturbed conditions is shown in Figure 6: under nominal dynamics (Figure 6a) the policy executes a coherent straight-line transit, while under 20% model uncertainty (Figure 7) the trajectory deviates substantially from the intended path. The ITAE under TC05 is approximately 5× larger than TC04, because the parametric perturbations cause torque prediction errors that accumulate and compound continuously over the full episode duration, producing large late-episode deviations weighted heavily by the ITAE time penalty.

The torque magnitude output under perturbed conditions, shown in Figure 7, exhibits high-frequency oscillation that is entirely absent under nominal dynamics, confirming that the L-module’s implicit representation of the plant input-output relationship becomes inaccurate when the effective dynamics shift. The UHRC’s recurrent carry provides no mechanism for online parameter adaptation, the carry states encode compressed representations of past observations and subgoals, not explicit physical model parameters.




Figure 5: Classical controller trajectory under parametric uncertainty (TC05, Success: 100%, RMSE: 0.707 m).

Prediction errors therefore compound irreversibly across successive timesteps with no corrective feedback equivalent to the PID integral term.

The 13.3% residual success rate, higher than the 0% recorded in TC04, reflects the nature of the failure geometry rather than improved robustness.

4 Discussion

The results across TC01–TC05 collectively reveal a clear and internally consistent picture. The UHRC’s dominant failure mode is not a failure of spatial reasoning but a failure of policy stability, the vehicle diverges from its intended path due to compounding torque prediction errors, irrespective of whether obstacles are present. This distinction is critical for interpreting the comparative results fairly. The classical controller’s failures, by contrast, are architecturally deterministic: the supervisory-regulatory interface introduces a replanning latency that is structurally irreducible without redesigning the pipeline. The two architectures therefore fail in qualitatively different ways that reveal complementary design boundaries rather than a simple performance ranking.

4.1 Policy Instability as the Dominant Failure Mode

Across TC01, TC02, and TC03, success rates of 6.7%, 3.3%, and 3.3% respectively confirm that the UHRC consistently fails to complete goal-directed navigation under nominal conditions. The mean RMSE values of 67.040 m, 45.476 m, and 48.413 m are far in excess of the geometric scale of each task (start-to-goal distances of approximately 10–18 m), confirming that the dominant failure is trajectory divergence rather than collision. This divergence originates from a distributional compounding problem, the behavioural cloning training objective minimises one-step prediction error, but at inference the model operates in a closed loop where every prediction error perturbs the state, causing the next




Figure 6: TC05 UHRC trajectory comparison between nominal and perturbed plant dynamics.

observation to drift progressively further from the training distribution Le Mero et al. 2022; Block, Foster, and Misra 2024.

The absence of an integral correction mechanism compounds this effect. The cascaded PID controller accumulates tracking error in the integrator state of each loop, providing a continuous restoring force that grows in proportion to the accumulated deviation. The UHRC has no equivalent mechanism. The carry state encodes a compressed representation of recent observations, but provides no signal that distinguishes a large persistent error from a sequence of transient deviations. Minor torque prediction errors therefore accumulate without bound until the vehicle enters states the policy has never encountered. The two-layer transformer in the L-module then operates in a regime where the input distribution has shifted entirely outside the training manifold, and the resulting torque commands drive further instability in a positive feedback loop.

The short-range successes in TC01 (2 of 30 trials with goal distances below approximately 7 m) confirm that the learned spatial reasoning is functionally correct when the episode terminates before compounding errors dominate. This is further supported by the qualitative trajectory evidence in TC02 and TC03, where the UHRC’s approach trajectories are smooth and geometrically appropriate the H-module correctly identifies the corridor passage and the gap as navigable subgoals but the vehicle fails to maintain stability long enough to execute them. The conclusion is that the spatial planning component of the UHRC is working as intended, but the action generation component lacks the error correction properties needed for stable long-horizon closed-loop operation.

4.2 Information Bottleneck vs Policy Stability: TC02 and TC03

The theoretical motivation for the UHRC is the information bottleneck at the classical controller’s supervisory-regulatory interface: the A* planner transmits only a waypoint sequence to the low-level PID, discarding all contextual information about obstacle prox-




Figure 7: UHRC total torque magnitude under 80% model certainty (TC05).

imity and local geometry Kolchinsky, Tracey, and Wolpert 2019; Abdelmaksoud, Mailah, and Abdallah 2020. The TC02 and TC03 results provide partial empirical evidence for this hypothesis. In both cases the classical controller exhibits lateral oscillation around the planned path near obstacle surfaces, attributable to the pure-pursuit tracker consuming only positional waypoints and having no representation of adjacent obstacle clearance. The UHRC’s single successful trajectories in these scenarios show qualitatively smoother clearance geometry, consistent with the unified observation vector providing the L-module with direct access to LiDAR returns during torque computation.

However, this architectural advantage is entirely negated by policy instability in 97% of trials. The information bottleneck hypothesis predicts that the UHRC should outperform the classical controller in constrained geometric environments. The results show instead that the UHRC cannot demonstrate this advantage because it fails to remain stable long enough for the spatial reasoning benefit to manifest. This does not invalidate the hypothesis, it establishes that resolving policy instability is a necessary precondition for testing it rigorously.

4.3 Architectural Failure Modes in TC04

TC04 is the only scenario in which both controllers record 0% success, yet the causal mechanisms are entirely distinct. The classical controller fails because the fixed supervisory-regulatory interface cannot respond to dynamic obstacles faster than the A* replanning cycle permits. The 10 Hz replanning rate and the ICP keyframe creation threshold of approximately one keyframe per 0.8 m of travel introduce a latency window during which the prior waypoint sequence remains active. This latency is not a consequence of poor tuning but of the fixed information pathway: the low-level PID has no access to sensor data and therefore cannot initiate avoidance until a revised waypoint sequence arrives from the supervisory layer S. Li et al. 2023; Zhu, Ferrari-Trecate, and Geroliminis 2024.

The UHRC fails for the unrelated reason of pre-existing policy instability, but the architectural analysis of its behaviour during coherent trials is informative. In those

trials where the vehicle remained on a stable path until $t = 3.0$ s, the LiDAR observation was updated within the same 10 ms inference cycle as the obstacle revelation, and the observation embedding reflected the new obstacle geometry immediately in the next H-L reasoning cycle. This is structurally impossible for the classical controller regardless of tuning. The UHRC-PID hybrid architecture in which the UHRC generates attitude references consumed by the inner PID attitude loops achieved a 20% success rate on TC04 in supplementary evaluation, providing direct empirical evidence that the UHRC’s reactive spatial reasoning is functionally present when the action generation instability is decoupled from the trajectory generation task. This result motivates the hybrid architecture as the more immediately viable path toward deployment.

4.4 Distributional Fragility vs Integral Robustness: TC05

TC05 most directly characterises the fundamental trade-off between the two architectures. The classical controller’s 100% success rate under simultaneous perturbations of mass ($\pm 20\%$), inertia ($\pm 30\%$), drag ($\pm 50\%$), and thrust coefficient ($\pm 15\%$) demonstrates that integral feedback is a principled and highly effective mechanism for parametric robustness. The deviations in effective plant dynamics are treated as persistent disturbances by the PID integrators and compensated continuously, without any requirement for the controller to have knowledge of the perturbation magnitude or structure. This robustness is architecturally guaranteed, it follows from the stability analysis of the cascaded PID structure under the small-signal linearisation used for gain selection, and does not depend on the specific perturbation values applied.

The UHRC’s 13.3% success rate under TC05 and its elevated ITAE of 73,011.560 m s² approximately $5\times$ larger than the next worst scenario confirm that parametric uncertainty is the most damaging distributional shift the policy encounters. Behavioural cloning on a single nominal dynamic model produces a policy whose implicit representation of the plant input-output relationship is calibrated exclusively to the nominal parameter set Le Mero et al. 2022. When the effective dynamics deviate, the L-module’s torque predictions become systematically biased, and the recurrent carry which encodes compressed representations of past observations and subgoals rather than explicit physical model parameters provides no mechanism for detecting and compensating the shift. The high-frequency torque oscillation observed in Figure 7 is the behavioural signature of this mismatch. The model generates rapid corrective torques in response to observation errors produced by the biased dynamics, which themselves produce further observation errors in a destabilising cycle.

The comparative results across TC01–TC05 therefore establish a clear performance-robustness boundary. The UHRC offers reactive spatial reasoning, fixed and bounded inference cost independent of mission duration, and a unified architecture that eliminates manual interface design. The classical controller provides formal stability guarantees, interpretable failure modes, integral disturbance rejection, and robustness to significant parametric variation without retraining. The UHRC’s advantages in spatial reactivity are currently negated by policy instability attributable to the limitations of offline behavioural cloning from a single nominal expert. Closing this gap is a question of training methodology rather than architecture.

5 Conclusion

This paper presented the Unified Hierarchical Reasoning Controller (UHRC), the first systematic application of the Hierarchical Reasoning Model (HRM) architecture to a continuous nonlinear closed-loop control system. The UHRC was implemented as a two-level recurrent transformer in which a high-level module generates body-frame velocity subgoals and a low-level module translates those subgoals into body wrench commands, trained via behavioural cloning on 500 A*-guided expert episodes. A persistent 16-slot rolling carry buffer, RoPE positional encoding, and a gated subgoal injection mechanism were introduced as principled adaptations of the discrete HRM architecture to the requirements of continuous real-time control.

The comparative evaluation against a classical SLAM-augmented cascaded PID baseline across five structured test cases produced principal findings. The HRM architecture can learn a functional continuous control policy. The policy instability was attributable to behavioural cloning’s compounding error limitation is the primary obstacle to reliable deployment. The two architectures fail in qualitatively distinct ways that reveal complementary design boundaries. Fourth, the performance-robustness trade-off between unified learned hierarchy and classical explicit decomposition is architectural in origin.

The primary direction for future work is domain randomisation during training, which directly addresses distributional fragility without architectural modification. Reinforcement learning fine-tuning from the behaviourally cloned initialisation would address compounding error instability by allowing the policy to discover corrective strategies absent from the offline expert. Beyond these immediate extensions, the UHRC framework provides a principled foundation for investigating whether unified hierarchical reasoning can match or exceed the performance of explicit architectural decomposition once the training methodology limitations are resolved a question this investigation has framed precisely but not yet answered.

References

- Abdelhay, Sabir and Alia Zakriti (2019). “Modeling of a Quadcopter Trajectory Tracking System Using PID Controller”. en. In: *Procedia Manufacturing* 32, pp. 564–571. ISSN: 23519789. DOI: [10.1016/j.promfg.2019.02.253](https://doi.org/10.1016/j.promfg.2019.02.253). URL: <https://linkinghub.elsevier.com/retrieve/pii/S2351978919302884> (visited on 10/21/2025).
- Abdelmaksoud, Sherif I., Musa Mailah, and Ayman M. Abdallah (2020). “Control Strategies and Novel Techniques for Autonomous Rotorcraft Unmanned Aerial Vehicles: A Review”. In: *IEEE Access* 8, pp. 195142–195169. ISSN: 2169-3536. DOI: [10.1109/ACCESS.2020.3031326](https://doi.org/10.1109/ACCESS.2020.3031326). URL: <https://ieeexplore.ieee.org/document/9224621/> (visited on 03/24/2026).
- Agvei, Klinsmann, Pouria Sarhadi, and Daniel Polani (July 2025). *Deep Reinforcement Learning in Applied Control: Challenges, Analysis, and Insights*. arXiv:2507.08196 [eess]. DOI: [10.48550/arXiv.2507.08196](https://doi.org/10.48550/arXiv.2507.08196). URL: <http://arxiv.org/abs/2507.08196> (visited on 11/18/2025).
- AlMahamid, Fadi and Katarina Grolinger (Sept. 2021). “Reinforcement Learning Algorithms: An Overview and Classification”. In: *2021 IEEE Canadian Conference on Electrical and Computer Engineering (CCECE)*. ON, Canada: IEEE, pp. 1–7. ISBN: 978-1-6654-4864-2. DOI: [10.1109/CCECE53047.2021.9569056](https://doi.org/10.1109/CCECE53047.2021.9569056). URL: <https://ieeexplore.ieee.org/document/9569056/> (visited on 09/22/2025).

- Ameli, Sina and Olugbenga Moses Anubi (Aug. 2021). *Robust Control for a Class of Non-linearly Coupled Hierarchical Systems with Actuator Faults*. arXiv:2106.14944 [math]. DOI: [10.48550/arXiv.2106.14944](https://doi.org/10.48550/arXiv.2106.14944). URL: <http://arxiv.org/abs/2106.14944> (visited on 09/14/2025).
- Badre, David and Derek Evan Nee (Feb. 2018). “Frontal Cortex and the Hierarchical Control of Behavior”. en. In: *Trends in Cognitive Sciences* 22.2, pp. 170–188. ISSN: 13646613. DOI: [10.1016/j.tics.2017.11.005](https://doi.org/10.1016/j.tics.2017.11.005). URL: <https://linkinghub.elsevier.com/retrieve/pii/S1364661317302450> (visited on 09/12/2025).
- Baheri, Ali and Cecilia O. Alm (May 2025). *Hierarchical Neuro-Symbolic Decision Transformer*. arXiv:2503.07148 [cs]. DOI: [10.48550/arXiv.2503.07148](https://doi.org/10.48550/arXiv.2503.07148). URL: <http://arxiv.org/abs/2503.07148> (visited on 09/17/2025).
- Bahnasawi, A.A., A.S. Al-Fuhaid, and M.S. Mahmoud (1990). “Decentralised and hierarchical control of interconnected uncertain systems”. en. In: *IEE Proceedings D Control Theory and Applications* 137.5, p. 311. ISSN: 01437054. DOI: [10.1049/ip-d.1990.0039](https://doi.org/10.1049/ip-d.1990.0039). URL: <https://digital-library.theiet.org/content/journals/10.1049/ip-d.1990.0039> (visited on 09/13/2025).
- Block, Adam, Dylan Foster, and Dipendra Misra (2024). “Is Behavior Cloning All You Need? Understanding Horizon in Imitation Learning”. In: *Advances in Neural Information Processing Systems 37*. Vancouver, BC, Canada: Neural Information Processing Systems Foundation, Inc. (NeurIPS), pp. 120602–120666. ISBN: 979-8-3313-1438-5. DOI: [10.52202/079017-3834](https://doi.org/10.52202/079017-3834). URL: <http://www.proceedings.com/079017-3834.html> (visited on 03/24/2026).
- Borri, Alessandro, Giordano Pola, and Maria Domenica Di Benedetto (Mar. 2019). “Design of Symbolic Controllers for Networked Control Systems”. In: *IEEE Transactions on Automatic Control* 64.3, pp. 1034–1046. ISSN: 0018-9286, 1558-2523, 2334-3303. DOI: [10.1109/TAC.2018.2833630](https://doi.org/10.1109/TAC.2018.2833630). URL: <https://ieeexplore.ieee.org/document/8355500/> (visited on 10/10/2025).
- Bouabdallah, S., A. Noth, and R. Siegwart (2004). “PID vs LQ control techniques applied to an indoor micro quadrotor”. In: *2004 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS) (IEEE Cat. No.04CH37566)*. Sendai, Japan: IEEE, 2451–2456 vol.3. ISBN: 978-0-7803-8463-7. DOI: [10.1109/IROS.2004.1389776](https://doi.org/10.1109/IROS.2004.1389776). URL: <https://ieeexplore.ieee.org/document/1389776/> (visited on 10/12/2025).
- Bouabdallah, Samir (2007). “Design and control of quadrotors with application to autonomous flying”. en. PhD thesis. Lausanne, EPFL. DOI: [10.5075/EPFL-THESIS-3727](https://doi.org/10.5075/EPFL-THESIS-3727). URL: <https://infoscience.epfl.ch/handle/20.500.14299/235984> (visited on 10/27/2025).
- Chen, Weifeng et al. (Nov. 2022). “SLAM Overview: From Single Sensor to Heterogeneous Fusion”. en. In: *Remote Sensing* 14.23, p. 6033. ISSN: 2072-4292. DOI: [10.3390/rs14236033](https://doi.org/10.3390/rs14236033). URL: <https://www.mdpi.com/2072-4292/14/23/6033> (visited on 03/24/2026).
- Dai, Hongkai et al. (Sept. 2021). *Lyapunov-stable neural-network control*. arXiv:2109.14152 [cs]. DOI: [10.48550/arXiv.2109.14152](https://doi.org/10.48550/arXiv.2109.14152). URL: <http://arxiv.org/abs/2109.14152> (visited on 10/11/2025).
- Daoutidis, Prodromos, Larry Megan, and Wentao Tang (Oct. 2023). “The future of control of process systems”. en. In: *Computers & Chemical Engineering* 178, p. 108365. ISSN: 00981354. DOI: [10.1016/j.compchemeng.2023.108365](https://doi.org/10.1016/j.compchemeng.2023.108365). URL: <https://linkinghub.elsevier.com/retrieve/pii/S0098135423002351> (visited on 09/20/2025).

- Daoutidis, Prodromos, Wentao Tang, and Andrew Allman (Oct. 2019). “Decomposition of control and optimization problems by network structure: Concepts, methods, and inspirations from biology”. en. In: *AIChE Journal* 65.10, e16708. ISSN: 0001-1541, 1547-5905. DOI: [10.1002/aic.16708](https://doi.org/10.1002/aic.16708). URL: <https://aiche.onlinelibrary.wiley.com/doi/10.1002/aic.16708> (visited on 09/14/2025).
- Dave, Akshay J. et al. (Mar. 2023). “Design of a Supervisory Control System for Autonomous Operation of Advanced Reactors”. In: *Annals of Nuclear Energy* 182. arXiv:2209.04334 [eess], p. 109593. ISSN: 03064549. DOI: [10.1016/j.anucene.2022.109593](https://doi.org/10.1016/j.anucene.2022.109593). URL: <http://arxiv.org/abs/2209.04334> (visited on 09/15/2025).
- Feng, Sicheng et al. (Sept. 2025). *Efficient Reasoning Models: A Survey*. arXiv:2504.10903 [cs]. DOI: [10.48550/arXiv.2504.10903](https://doi.org/10.48550/arXiv.2504.10903). URL: <http://arxiv.org/abs/2504.10903> (visited on 10/10/2025).
- Geiping, Jonas et al. (Feb. 2025). *Scaling up Test-Time Compute with Latent Reasoning: A Recurrent Depth Approach*. arXiv:2502.05171 [cs]. DOI: [10.48550/arXiv.2502.05171](https://doi.org/10.48550/arXiv.2502.05171). URL: <http://arxiv.org/abs/2502.05171> (visited on 09/12/2025).
- Goodwin, Graham Clifford, Stefan F. Graebe, and Mario E. Salgado (2001). *Control System Design*. en. Google-Books-ID: 7dNSAAAAMAAJ. Prentice Hall. ISBN: 978-0-13-958653-8.
- Gupta, Abhishek and Xavier Fernando (Mar. 2022). “Simultaneous Localization and Mapping (SLAM) and Data Fusion in Unmanned Aerial Vehicles: Recent Advances and Challenges”. en. In: *Drones* 6.4, p. 85. ISSN: 2504-446X. DOI: [10.3390/drones6040085](https://doi.org/10.3390/drones6040085). URL: <https://www.mdpi.com/2504-446X/6/4/85> (visited on 03/24/2026).
- Hoffmann, Gabriel et al. (Aug. 2007). “Quadrotor Helicopter Flight Dynamics and Control: Theory and Experiment”. en. In: *AIAA Guidance, Navigation and Control Conference and Exhibit*. Hilton Head, South Carolina: American Institute of Aeronautics and Astronautics. ISBN: 978-1-62410-015-4. DOI: [10.2514/6.2007-6461](https://doi.org/10.2514/6.2007-6461). URL: <https://arc.aiaa.org/doi/10.2514/6.2007-6461> (visited on 10/12/2025).
- Huang, Jie and Kevin Chen-Chuan Chang (May 2023). *Towards Reasoning in Large Language Models: A Survey*. arXiv:2212.10403 [cs]. DOI: [10.48550/arXiv.2212.10403](https://doi.org/10.48550/arXiv.2212.10403). URL: <http://arxiv.org/abs/2212.10403> (visited on 09/12/2025).
- Huang, Yaofu et al. (Oct. 2020). “Path Tracking Based on Improved Pure Pursuit Model and PID”. In: *2020 IEEE 2nd International Conference on Civil Aviation Safety and Information Technology (ICCASIT)*. Weihai, China: IEEE, pp. 359–364. ISBN: 978-1-7281-9948-1. DOI: [10.1109/ICCASIT50869.2020.9368694](https://doi.org/10.1109/ICCASIT50869.2020.9368694). URL: <https://ieeexplore.ieee.org/document/9368694/> (visited on 03/24/2026).
- Huntenburg, Julia M., Pierre-Louis Bazin, and Daniel S. Margulies (Jan. 2018). “Large-Scale Gradients in Human Cortical Organization”. English. In: *Trends in Cognitive Sciences* 22.1, pp. 21–31. ISSN: 1364-6613, 1879-307X. DOI: [10.1016/j.tics.2017.11.002](https://doi.org/10.1016/j.tics.2017.11.002). URL: [https://www.cell.com/trends/cognitive-sciences/abstract/S1364-6613\(17\)30240-1](https://www.cell.com/trends/cognitive-sciences/abstract/S1364-6613(17)30240-1) (visited on 09/10/2025).
- Iqbal, Jamshed et al. (Dec. 2017). “Nonlinear control systems - A brief overview of historical and recent advances”. In: *Nonlinear Engineering* 6.4. ISSN: 2192-8010, 2192-8029. DOI: [10.1515/nleng-2016-0077](https://doi.org/10.1515/nleng-2016-0077). URL: <https://www.degruyter.com/document/doi/10.1515/nleng-2016-0077/html> (visited on 09/20/2025).
- Jin, Jiajie et al. (July 2025). *Decoupled Planning and Execution: A Hierarchical Reasoning Framework for Deep Search*. arXiv:2507.02652 [cs]. DOI: [10.48550/arXiv.2507.02652](https://doi.org/10.48550/arXiv.2507.02652). URL: <http://arxiv.org/abs/2507.02652> (visited on 09/12/2025).

- Kiran, B Ravi et al. (June 2022). “Deep Reinforcement Learning for Autonomous Driving: A Survey”. In: *IEEE Transactions on Intelligent Transportation Systems* 23.6, pp. 4909–4926. ISSN: 1524-9050, 1558-0016. DOI: [10.1109/TITS.2021.3054625](https://doi.org/10.1109/TITS.2021.3054625). URL: <https://ieeexplore.ieee.org/document/9351818/> (visited on 09/22/2025).
- Kolchinsky, Artemy, Brendan D. Tracey, and David H. Wolpert (Nov. 2019). “Nonlinear Information Bottleneck”. en. In: *Entropy* 21.12, p. 1181. ISSN: 1099-4300. DOI: [10.3390/e21121181](https://doi.org/10.3390/e21121181). URL: <https://www.mdpi.com/1099-4300/21/12/1181> (visited on 03/25/2026).
- LaValle, Steven M. (2006). *Planning algorithms*. eng. New York: Cambridge University Press. ISBN: 978-0-521-86205-9.
- Le Mero, Luc et al. (Sept. 2022). “A Survey on Imitation Learning Techniques for End-to-End Autonomous Vehicles”. In: *IEEE Transactions on Intelligent Transportation Systems* 23.9, pp. 14128–14147. ISSN: 1524-9050, 1558-0016. DOI: [10.1109/TITS.2022.3144867](https://doi.org/10.1109/TITS.2022.3144867). URL: <https://ieeexplore.ieee.org/document/9700770/> (visited on 03/24/2026).
- Levine, Sergey et al. (Apr. 2016). *End-to-End Training of Deep Visuomotor Policies*. arXiv:1504.00702 [cs]. DOI: [10.48550/arXiv.1504.00702](https://doi.org/10.48550/arXiv.1504.00702). URL: <http://arxiv.org/abs/1504.00702> (visited on 09/15/2025).
- Li, Hengli et al. (May 2025). *Seek in the Dark: Reasoning via Test-Time Instance-Level Policy Gradient in Latent Space*. arXiv:2505.13308 [cs]. DOI: [10.48550/arXiv.2505.13308](https://doi.org/10.48550/arXiv.2505.13308). URL: <http://arxiv.org/abs/2505.13308> (visited on 09/12/2025).
- Li, Sijia et al. (June 2023). “Hierarchical Control for Microgrids: A Survey on Classical and Machine Learning-Based Methods”. en. In: *Sustainability* 15.11, p. 8952. ISSN: 2071-1050. DOI: [10.3390/su15118952](https://doi.org/10.3390/su15118952). URL: <https://www.mdpi.com/2071-1050/15/11/8952> (visited on 09/12/2025).
- Li, Zhong-Zhi et al. (June 2025). *From System 1 to System 2: A Survey of Reasoning Large Language Models*. arXiv:2502.17419 [cs]. DOI: [10.48550/arXiv.2502.17419](https://doi.org/10.48550/arXiv.2502.17419). URL: <http://arxiv.org/abs/2502.17419> (visited on 09/24/2025).
- Liang, Xiao et al. (Apr. 2018). “Nonlinear Hierarchical Control for Unmanned Quadrotor Transportation Systems”. In: *IEEE Transactions on Industrial Electronics* 65.4, pp. 3395–3405. ISSN: 0278-0046, 1557-9948. DOI: [10.1109/TIE.2017.2752139](https://doi.org/10.1109/TIE.2017.2752139). URL: <http://ieeexplore.ieee.org/document/8038248/> (visited on 09/12/2025).
- Liu, Mingxing, Yang Tan, and Vincent Padois (Jan. 2016). “Generalized hierarchical control”. en. In: *Autonomous Robots* 40.1, pp. 17–31. ISSN: 0929-5593, 1573-7527. DOI: [10.1007/s10514-015-9436-1](https://doi.org/10.1007/s10514-015-9436-1). URL: <http://link.springer.com/10.1007/s10514-015-9436-1> (visited on 09/12/2025).
- Liu, Wei, Zheng Zheng, and Kaiyuan Cai (June 2013). “Adaptive path planning for unmanned aerial vehicles based on bi-level programming and variable planning time interval”. en. In: *Chinese Journal of Aeronautics* 26.3, pp. 646–660. ISSN: 10009361. DOI: [10.1016/j.cja.2013.04.041](https://doi.org/10.1016/j.cja.2013.04.041). URL: <https://linkinghub.elsevier.com/retrieve/pii/S1000936113000927> (visited on 09/13/2025).
- Liu, Zhitao et al. (2015). “New Trends in Nonlinear Control Systems and Applications”. en. In: *Abstract and Applied Analysis* 2015, pp. 1–2. ISSN: 1085-3375, 1687-0409. DOI: [10.1155/2015/637632](https://doi.org/10.1155/2015/637632). URL: <http://www.hindawi.com/journals/aaa/2015/637632/> (visited on 09/20/2025).
- Lopez-Sanchez, Ivan and Javier Moreno-Valenzuela (2023). “PID control of quadrotor UAVs: A survey”. en. In: *Annual Reviews in Control* 56, p. 100900. ISSN: 13675788.

- DOI: [10.1016/j.arcontrol.2023.100900](https://doi.org/10.1016/j.arcontrol.2023.100900). URL: <https://linkinghub.elsevier.com/retrieve/pii/S1367578823000640> (visited on 03/24/2026).
- Ly, Abdoulaye O. and Moulay Akhloufi (June 2021). “Learning to Drive by Imitation: An Overview of Deep Behavior Cloning Methods”. In: *IEEE Transactions on Intelligent Vehicles* 6.2, pp. 195–209. ISSN: 2379-8904, 2379-8858. DOI: [10.1109/TIV.2020.3002505](https://doi.org/10.1109/TIV.2020.3002505). URL: <https://ieeexplore.ieee.org/document/9117169/> (visited on 03/24/2026).
- Mahmoud, Tarek A. et al. (Apr. 2024). “Fractional-order fuzzy sliding mode control of uncertain nonlinear MIMO systems using fractional-order reinforcement learning”. en. In: *Complex & Intelligent Systems* 10.2, pp. 3057–3085. ISSN: 2199-4536, 2198-6053. DOI: [10.1007/s40747-023-01309-8](https://doi.org/10.1007/s40747-023-01309-8). URL: <https://link.springer.com/10.1007/s40747-023-01309-8> (visited on 09/22/2025).
- Mesarovic, M.D., D. Macko, and Y. Takahara (Mar. 1970). “Two coordination principles and their application in large scale systems control”. en. In: *Automatica* 6.2, pp. 261–270. ISSN: 00051098. DOI: [10.1016/0005-1098\(70\)90097-X](https://doi.org/10.1016/0005-1098(70)90097-X). URL: <https://linkinghub.elsevier.com/retrieve/pii/000510987090097X> (visited on 09/12/2025).
- Milijas, Robert et al. (June 2021). “A Comparison of LiDAR-based SLAM Systems for Control of Unmanned Aerial Vehicles”. In: *2021 International Conference on Unmanned Aircraft Systems (ICUAS)*. Athens, Greece: IEEE, pp. 1148–1154. ISBN: 978-1-6654-1535-4. DOI: [10.1109/ICUAS51884.2021.9476802](https://doi.org/10.1109/ICUAS51884.2021.9476802). URL: <https://ieeexplore.ieee.org/document/9476802/> (visited on 01/10/2026).
- Munguia, Rodrigo, Juan-Carlos Trujillo, and Antoni Grau (Aug. 2025). “UAV Navigation Using EKF-MonoSLAM Aided by Range-to-Base Measurements”. en. In: *Drones* 9.8, p. 570. ISSN: 2504-446X. DOI: [10.3390/drones9080570](https://doi.org/10.3390/drones9080570). URL: <https://www.mdpi.com/2504-446X/9/8/570> (visited on 01/10/2026).
- Murray, John D. et al. (Dec. 2014). “A hierarchy of intrinsic timescales across primate cortex”. en. In: *Nature Neuroscience* 17.12, pp. 1661–1663. ISSN: 1546-1726. DOI: [10.1038/nn.3862](https://doi.org/10.1038/nn.3862). URL: <https://www.nature.com/articles/nn.3862> (visited on 09/10/2025).
- Nagaty, Amr et al. (Apr. 2013). “Control and Navigation Framework for Quadrotor Helicopters”. en. In: *Journal of Intelligent & Robotic Systems* 70.1-4, pp. 1–12. ISSN: 0921-0296, 1573-0409. DOI: [10.1007/s10846-012-9789-z](https://doi.org/10.1007/s10846-012-9789-z). URL: <http://link.springer.com/10.1007/s10846-012-9789-z> (visited on 10/13/2025).
- Nguyen, D. H., M. H. Lowenberg, and S. A. Neild (Apr. 2021). “Identifying limits of linear control design validity in nonlinear systems: a continuation-based approach”. en. In: *Nonlinear Dynamics* 104.2, pp. 901–921. ISSN: 1573-269X. DOI: [10.1007/s11071-021-06341-2](https://doi.org/10.1007/s11071-021-06341-2). URL: <https://doi.org/10.1007/s11071-021-06341-2> (visited on 09/14/2025).
- Qin, Hongwei et al. (Mar. 2023). “Review of Autonomous Path Planning Algorithms for Mobile Robots”. en. In: *Drones* 7.3, p. 211. ISSN: 2504-446X. DOI: [10.3390/drones7030211](https://doi.org/10.3390/drones7030211). URL: <https://www.mdpi.com/2504-446X/7/3/211> (visited on 09/17/2025).
- Ren, Jiao et al. (Aug. 2016). “Cascade PID controller for quadrotor”. In: *2016 IEEE International Conference on Information and Automation (ICIA)*. Ningbo, China: IEEE, pp. 120–124. ISBN: 978-1-5090-4102-2. DOI: [10.1109/ICInfA.2016.7831807](https://doi.org/10.1109/ICInfA.2016.7831807). URL: <https://ieeexplore.ieee.org/document/7831807/> (visited on 10/21/2025).
- Ren, Tianyu and Houtan Jebelli (May 2025). “Efficient 3D robotic mapping and navigation method in complex construction environments”. en. In: *Computer-Aided Civil and*

- Infrastructure Engineering* 40.12, pp. 1580–1605. ISSN: 10939687. DOI: [10.1111/mice.13353](https://doi.org/10.1111/mice.13353). URL: <https://linkinghub.elsevier.com/retrieve/pii/S1093968726000137> (visited on 03/24/2026).
- Roberts, Peter D. and Victor M. Becerra (Nov. 2001). “Optimal control of a class of discrete–continuous non-linear systems—decomposition and hierarchical structure”. en. In: *Automatica* 37.11, pp. 1757–1769. ISSN: 00051098. DOI: [10.1016/S0005-1098\(01\)00141-8](https://doi.org/10.1016/S0005-1098(01)00141-8). URL: <https://linkinghub.elsevier.com/retrieve/pii/S0005109801001418> (visited on 09/13/2025).
- Rojas, Luis, Víctor Yepes, and José Garcia (Mar. 2025). “Complex Dynamics and Intelligent Control: Advances, Challenges, and Applications in Mining and Industrial Processes”. en. In: *Mathematics* 13.6, p. 961. ISSN: 2227-7390. DOI: [10.3390/math13060961](https://doi.org/10.3390/math13060961). URL: <https://www.mdpi.com/2227-7390/13/6/961> (visited on 09/14/2025).
- Sadati, N. and M. H. Ramezani (June 2008). “Hierarchical optimal control of nonlinear systems; An application to a benchmark CSTR problem”. In: *2008 3rd IEEE Conference on Industrial Electronics and Applications*. Singapore: IEEE, pp. 455–460. ISBN: 978-1-4244-1717-9. DOI: [10.1109/ICIEA.2008.4582557](https://doi.org/10.1109/ICIEA.2008.4582557). URL: <https://ieeexplore.ieee.org/document/4582557/> (visited on 09/12/2025).
- Sadati, Nasser, Mehdi Rahmani, and Mehrdad Saif (Mar. 2015). “Two-Level Robust Optimal Control of Large-Scale Nonlinear Systems”. In: *IEEE Systems Journal* 9.1, pp. 242–251. ISSN: 1932-8184, 1937-9234, 2373-7816. DOI: [10.1109/JSYST.2013.2287772](https://doi.org/10.1109/JSYST.2013.2287772). URL: <http://ieeexplore.ieee.org/document/6671992/> (visited on 09/13/2025).
- Sanders, C.P. et al. (1998). “Hierarchical control of small autonomous helicopters”. In: *Proceedings of the 37th IEEE Conference on Decision and Control (Cat. No.98CH36171)*. Vol. 4. Tampa, FL, USA: IEEE, pp. 3629–3634. ISBN: 978-0-7803-4394-8. DOI: [10.1109/CDC.1998.761744](https://doi.org/10.1109/CDC.1998.761744). URL: <http://ieeexplore.ieee.org/document/761744/> (visited on 09/13/2025).
- Saviolo, Alessandro and Giuseppe Loianno (2023). “Learning quadrotor dynamics for precise, safe, and agile flight control”. en. In: *Annual Reviews in Control* 55, pp. 45–60. ISSN: 13675788. DOI: [10.1016/j.arcontrol.2023.03.009](https://doi.org/10.1016/j.arcontrol.2023.03.009). URL: <https://linkinghub.elsevier.com/retrieve/pii/S1367578823000135> (visited on 10/12/2025).
- Scattolini, Riccardo and Patrizio Colaneri (2007). “Hierarchical model predictive control”. In: *2007 46th IEEE Conference on Decision and Control*. New Orleans, LA, USA: IEEE, pp. 4803–4808. ISBN: 978-1-4244-1497-0. DOI: [10.1109/CDC.2007.4434079](https://doi.org/10.1109/CDC.2007.4434079). URL: <http://ieeexplore.ieee.org/document/4434079/> (visited on 09/12/2025).
- Shakya, Ashish Kumar, Gopinatha Pillai, and Sohom Chakrabarty (Nov. 2023). “Reinforcement learning algorithms: A brief survey”. en. In: *Expert Systems with Applications* 231, p. 120495. ISSN: 09574174. DOI: [10.1016/j.eswa.2023.120495](https://doi.org/10.1016/j.eswa.2023.120495). URL: <https://linkinghub.elsevier.com/retrieve/pii/S0957417423009971> (visited on 09/22/2025).
- Shastri, Abhishek K., Mangal Kothari, and Abhishek Abhishek (Jan. 2018). “Development of Flight Dynamics Model of Quadrotor”. en. In: *2018 AIAA Guidance, Navigation, and Control Conference*. Kissimmee, Florida: American Institute of Aeronautics and Astronautics. ISBN: 978-1-62410-526-5. DOI: [10.2514/6.2018-1850](https://doi.org/10.2514/6.2018-1850). URL: <https://arc.aiaa.org/doi/10.2514/6.2018-1850> (visited on 10/12/2025).
- Siljak, Dragoslav D. (2013). *Decentralized Control of Complex Systems*. eng. OCLC: 1162446576. Newburyport: Dover Publications. ISBN: 978-0-486-29437-7.

- Singh, Madan G. and André Titli (1978). *Systems: decomposition, optimisation, and control*. Pergamon international library of science, technology, engineering, and social studies. Oxford ; New York: Pergamon Press. ISBN: 978-0-08-022150-2 978-0-08-023238-6.
- Slotine, J.-J. E. and Weiping Li (1991). *Applied nonlinear control*. Englewood Cliffs, N.J: Prentice Hall. ISBN: 978-0-13-040890-7.
- Tang, Hooi Hung and Nur Syazreen Ahmad (Dec. 2024). “Fuzzy logic approach for controlling uncertain and nonlinear systems: a comprehensive review of applications and advances”. en. In: *Systems Science & Control Engineering* 12.1, p. 2394429. ISSN: 2164-2583. DOI: [10.1080/21642583.2024.2394429](https://doi.org/10.1080/21642583.2024.2394429). URL: <https://www.tandfonline.com/doi/full/10.1080/21642583.2024.2394429> (visited on 09/21/2025).
- Torrente, Guillem et al. (Apr. 2021). “Data-Driven MPC for Quadrotors”. In: *IEEE Robotics and Automation Letters* 6.2, pp. 3769–3776. ISSN: 2377-3766, 2377-3774. DOI: [10.1109/LRA.2021.3061307](https://doi.org/10.1109/LRA.2021.3061307). URL: <https://ieeexplore.ieee.org/document/9361343/> (visited on 09/20/2025).
- Trad, Taha Yacine et al. (2024). “Real-Time Implementation of Quadrotor UAV Control System Based on a Deep Reinforcement Learning Approach”. en. In: *Computers, Materials & Continua* 81.3, pp. 4757–4786. ISSN: 1546-2226. DOI: [10.32604/cmc.2024.055634](https://doi.org/10.32604/cmc.2024.055634). URL: <https://www.techscience.com/cmc/v81n3/59024> (visited on 03/24/2026).
- Tran, Vu Phi et al. (June 2021). “Hybrid adaptive negative imaginary- neural-fuzzy control with model identification for a quadrotor”. en. In: *IFAC Journal of Systems and Control* 16, p. 100156. ISSN: 24686018. DOI: [10.1016/j.ifacsc.2021.100156](https://doi.org/10.1016/j.ifacsc.2021.100156). URL: <https://linkinghub.elsevier.com/retrieve/pii/S2468601821000110> (visited on 10/10/2025).
- Wang, Guan et al. (Aug. 2025). *Hierarchical Reasoning Model*. arXiv:2506.21734 [cs]. DOI: [10.48550/arXiv.2506.21734](https://doi.org/10.48550/arXiv.2506.21734). URL: <http://arxiv.org/abs/2506.21734> (visited on 09/10/2025).
- Wang, Xue-Fang, Jingjing Jiang, and Wen-Hua Chen (Jan. 2024). *A hierarchical control framework for autonomous decision-making systems: Integrating HMDP and MPC*. arXiv:2401.06833 [eess]. DOI: [10.48550/arXiv.2401.06833](https://doi.org/10.48550/arXiv.2401.06833). URL: <http://arxiv.org/abs/2401.06833> (visited on 09/12/2025).
- Wilson, I. D. (June 1979). “Foundations of hierarchical control”. en. In: *International Journal of Control* 29.6, pp. 899–933. ISSN: 0020-7179, 1366-5820. DOI: [10.1080/00207177908922740](https://doi.org/10.1080/00207177908922740). URL: <https://www.tandfonline.com/doi/full/10.1080/00207177908922740> (visited on 09/15/2025).
- Yu, Dongran et al. (June 2023). *A Survey on Neural-symbolic Learning Systems*. arXiv:2111.08164 [cs]. DOI: [10.48550/arXiv.2111.08164](https://doi.org/10.48550/arXiv.2111.08164). URL: <http://arxiv.org/abs/2111.08164> (visited on 09/12/2025).
- Zhang, Guangquan, Jie Lu, and Ya Gao (2015). *Multi-Level Decision Making: Models, Methods and Applications*. en. Vol. 82. Intelligent Systems Reference Library. Berlin, Heidelberg: Springer Berlin Heidelberg. ISBN: 978-3-662-46058-0 978-3-662-46059-7. DOI: [10.1007/978-3-662-46059-7](https://doi.org/10.1007/978-3-662-46059-7). URL: <https://link.springer.com/10.1007/978-3-662-46059-7> (visited on 09/13/2025).
- Zhou, Chengmin, Bingding Huang, and Pasi Fränti (Feb. 2021). *A review of motion planning algorithms for intelligent robotics*. arXiv:2102.02376 [cs]. DOI: [10.48550/arXiv.2102.02376](https://doi.org/10.48550/arXiv.2102.02376). URL: <http://arxiv.org/abs/2102.02376> (visited on 09/17/2025).

- Zhou, Zhao et al. (Mar. 2017). “Two-Level Hierarchical Model-Based Predictive Control for Large-Scale Urban Traffic Networks”. In: *IEEE Transactions on Control Systems Technology* 25.2, pp. 496–508. ISSN: 1063-6536, 1558-0865. DOI: [10.1109/TCST.2016.2572169](https://doi.org/10.1109/TCST.2016.2572169). URL: <http://ieeexplore.ieee.org/document/7488982/> (visited on 09/15/2025).
- Zhu, Pengbo, Giancarlo Ferrari-Trecate, and Nikolas Geroliminis (June 2024). *Hierarchical Control for Vehicle Repositioning in Autonomous Mobility on Demand Systems*. arXiv:2406.09609 [eess] version: 1. DOI: [10.48550/arXiv.2406.09609](https://doi.org/10.48550/arXiv.2406.09609). URL: <http://arxiv.org/abs/2406.09609> (visited on 09/12/2025).